

Using Machine Learning to Create Price Prediction for Tokyo Based Real Estate Company with Python

Part A: Data Processing and Exploratory Data Analysis (EDA)

By Shahin Karami

This Notebook is split into 2 steps. Step 1 will take Real Estate Data from Tokyo Japan (data is sourced from Kaggle) and prepare it for EDA, Correlation analysis, and eventually Regression analysis using Machine Learning. I will focus on removing unnecessary columns, duplicate values and missing values. Once the data is properly cleaned I move to Step 2, which is to conduct basic EDA with the help of `SweetViz`. I also remove outliers using the Median Absolute Deviation. By the conclusion of this notebook I will have prepped the data for all forms of analysis, and will have taken the first step in that analysis via EDA.

Step 1: Data Preprocessing

First I will prep the data for EDA.

1 - Imports and Data Import

The data used for this project comes [Kaggle](#). The Kaggle Dataset has real estate data from all 47 prefectures in Japan. For this project I use a subsample of the Tokyo prefecture data that only has 20,000 data points.

```
In [1]: import pickle
import matplotlib
from scipy.stats import median_abs_deviation

import pandas as pd
import numpy as np
import sweetviz as sv
import seaborn as sns

import matplotlib.pyplot as plt
import matplotlib.mlab as mlab
from matplotlib.pyplot import figure

plt.style.use('ggplot')
%matplotlib inline
```

```
In [2]: df = pd.read_csv('TokyoSub.csv')
```

2 - Removing Unessasary Columns

First I observe the dimensionality of the data, then I proceed to remove the columns that I will not be using for EDA.

```
In [3]: df.shape
```

```
Out[3]: (20000, 35)
```

```
In [4]: df = df.drop(['MunicipalityCode', 'Prefecture', 'Region', 'DistrictName', 'NearestStation', 'MinToNearestStation', 'AreaIsGreaterFlag', 'FrontageIsGreaterFlag', 'TotalFloorArea', 'TotalFloorAreaIsGreaterFlag', 'PrewarBuilding', 'Purpose', 'Classification', 'BuildingYear', 'CityPlanning', 'CoverageRatio', 'FloorAreaRatio', 'Period', 'Year', 'Quarter', 'Renovation'])
```

```
In [5]: df.head(5)
```

	Type	Municipality	TimeToNearestStation	TradePrice	FloorPlan	Area	LandShape	Frontage
0	Pre-owned Condominiums, etc.	Chiyoda Ward	4	400000000	1LDK	30	NaN	NaN
1	Pre-owned Condominiums, etc.	Chiyoda Ward	4	1300000000	3LDK	80	NaN	NaN
2	Residential Land(Land and Building)	Chiyoda Ward	2	4000000000	NaN	110	Trapezoidal Shaped	9.0
3	Residential Land(Land and Building)	Chiyoda Ward	1	1800000000	NaN	50	Rectangular Shaped	5.2
4	Pre-owned Condominiums, etc.	Chiyoda Ward	4	1000000000	2LDK	65	NaN	NaN

3 - Change Column's Name

I rename some columns to make the data easier to interpret. I change the name of the column based on the following:


- `TimeToNearestStation` : `MinuteToClosestStation`
- `TradePrice` : `Price`

```
In [6]: col_name = ['Type', 'Municipality', 'MinuteToClosestStation', 'Price', 'FloorPlan', 'Area', 'LandShape', 'Frontage', 'BuildingYear', 'Structure', 'Use', 'Direction', 'Renovation']

df.columns = col_name
df.head(5)
```

Out[6]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage
0	Pre-owned Condominiums, etc.	Chiyoda Ward	4	40000000	1LDK	30	NaN	NaN
1	Pre-owned Condominiums, etc.	Chiyoda Ward	4	130000000	3LDK	80	NaN	NaN
2	Residential Land(Land and Building)	Chiyoda Ward	2	400000000	NaN	110	Trapezoidal Shaped	9.
3	Residential Land(Land and Building)	Chiyoda Ward	1	180000000	NaN	50	Rectangular Shaped	5.
4	Pre-owned Condominiums, etc.	Chiyoda Ward	4	100000000	2LDK	65	NaN	NaN



4 - Dealing with Missing Values

I create a for loop to iterate among all columns to determine whether they have any `NaN` or not. This Loop provides the number of null values in every single column as well as the percentage of null values.

```
In [7]: for col in df.columns:
        number_null = df.loc[:, col].isnull().sum()
        perc_null = (number_null / df.shape[0]) * 100
        print('{} -- {} -- {}'.format(col, number_null, round(perc_null,3)))
```

```
Type -- 0 -- 0.0%
Municipality -- 0 -- 0.0%
MinuteToClosestStation -- 137 -- 0.685%
Price -- 0 -- 0.0%
FloorPlan -- 3153 -- 15.765%
Area -- 0 -- 0.0%
LandShape -- 16848 -- 84.24%
Frontage -- 17053 -- 85.265%
BuildingYear -- 710 -- 3.55%
Structure -- 488 -- 2.44%
Use -- 2208 -- 11.04%
Direction -- 16847 -- 84.235%
Renovation -- 4917 -- 24.585%
```

The result above shows we have a wide range of `NaN` present in the data. Some features contain a small amount of `NaN` while other contain a large amount of NaNs. Since I have a large dataset I have created a trivial breakdown of how I will go about it the `NaN` within each feature:

Numerical Features:

- between 0% and 30%, I impute with mean or median.

- More than 30%, I drop the rows.

Categorical Features:

- less than 5%, I drop the rows.
- between 5% and 30%, I impute with mode.
- More than 30%, create a new label as `Other` .

NOTE: *By no means do I claim this method to be the 'correct' method. However, for the sake of this project this method will suffice.*

I first focus on the categorical values. I impute `Renovation` , `FloorPlan` and `Use` with mode.

```
In [8]: mode_Renovation = df['Renovation'].mode()[0]
mode_Use = df['Use'].mode()[0]
mode_FloorPlan = df['FloorPlan'].mode()[0]

print(mode_Renovation)
print(mode_Use)
print(mode_FloorPlan)
```

```
Not yet
House
1K
```

```
In [9]: df['Renovation'] = df['Renovation'].fillna('Not yet')
df['Use'] = df['Use'].fillna('House')
df['FloorPlan'] = df['FloorPlan'].fillna('1K')
```

Then, I created a new label for missing values in `LandShape` and `Direction` .

```
In [10]: df['LandShape'] = df['LandShape'].fillna('Other LandShape')
df['Direction'] = df['Direction'].fillna('Other Direction')
```

Lastly, I drop any row with `NaNs` in the `Structure` columns

```
In [11]: df = df.dropna(subset = ['Structure'])
```

Now, I will shift my focus to the Numerical Features. I will impute `MinuteToClosestStation` and `BuildingYear` . However, before doing this, I must make sure about distribution shape of these columns to see whether they are right-skewed or left-skewed. It can be helpful when I want to decide choosing mean or median for imputing. Also, I will check the data type of these variables to confirm their numerical type.

```
In [12]: print(df['MinuteToClosestStation'].dtypes)
print(df['BuildingYear'].dtypes)
```

```
object
float64
```

`MinuteToClosestStation` is an object type. This means there are some non-numerical values in this column. This should be handled first. Below we will confirm the presence of non-numerical values and instruct python to coerce those instances as a `NaN` .

```
In [13]: df['MinuteToClosestStation'].unique()
```

```
Out[13]: array(['4', '2', '1', nan, '3', '5', '6', '0', '7', '11', '10', '8', '9',  
              '12', '30-60minutes', '13', '14', '16', '15', '19', '18', '21',  
              '2H-', '17'], dtype=object)
```

```
In [14]: df['MinuteToClosestStation'] = pd.to_numeric(df['MinuteToClosestStation'], errors='coerce')  
df['MinuteToClosestStation'].unique()
```

```
Out[14]: array([ 4.,  2.,  1., nan,  3.,  5.,  6.,  0.,  7., 11., 10.,  8.,  9.,  
              12., 13., 14., 16., 15., 19., 18., 21., 17.])
```

Now we can see its distribution to determine whether to use Mean or Median

```
In [15]: skewness_MTCS = round(df['MinuteToClosestStation'].skew(),3)  
print('Skewness : ' , skewness_MTCS)  
  
def approximately_zero(skewness, tolerance= 1):  
    if abs(skewness) <= tolerance:  
        print("The data is not heavily skewed, use the Mean. ")  
    else:  
        print("The data is heavily skewed, use the Median. ")  
    return  
  
approximately_zero(skewness_MTCS)
```

Skewness : 1.103

The data is heavily skewed, use the Median.

The result above tells me that I should opt to use the Median when I impute.

```
In [16]: df['MinuteToClosestStation'] = df['MinuteToClosestStation'].fillna(df['MinuteToClosestStation'].median())
```

Next I run the same process for `BuildingYear`

```
In [17]: skewness_BY = round(df['BuildingYear'].skew(),3)  
print('Skewness : ' , skewness_BY)  
  
approximately_zero(skewness_BY)
```

Skewness : -0.933

The data is not heavily skewed, use the Mean.

The result above tells me that I should opt to use the Mean when I impute.

```
In [18]: df['BuildingYear'] = df['BuildingYear'].fillna(df['BuildingYear'].mean()).round(0)
```

The last change i need to do is to drop the `NaNs` found in the `Frontage` columns

```
In [19]: df = df.dropna(subset = ['Frontage'])
```

Now that I have completed all imputions I conduct I final check in the updated data.

```
In [20]: for col in df.columns:  
    number_null = df.loc[:, col].isnull().sum()  
    perc_null = (number_null / df.shape[0]) * 100  
    print('{ } - { } - {}'.format(col, number_null, round(perc_null,3)))
```

```
Type - 0 - %0.0
Municipality - 0 - %0.0
MinuteToClosestStation - 0 - %0.0
Price - 0 - %0.0
FloorPlan - 0 - %0.0
Area - 0 - %0.0
LandShape - 0 - %0.0
Frontage - 0 - %0.0
BuildingYear - 0 - %0.0
Structure - 0 - %0.0
Use - 0 - %0.0
Direction - 0 - %0.0
Renovation - 0 - %0.0
```

5 - Removing any Duplicates Entries.

Now that all the `NaNs` have been handled, I can focus on removing any potential duplicated. I will first check to see if the data contains duplicates, then proceed accordingly.

```
In [21]: duplicate_rows = df.duplicated()

if duplicate_rows.any():
    print("The DataFrame has duplicate rows.")
else:
    print("The DataFrame does not have duplicate rows.")
```

The DataFrame has duplicate rows.

Since we have duplicates rows, lets find them and list them out in a DataFrame

```
In [22]: duplicate_rows = df[df.duplicated(keep = False)]
duplicate_rows = pd.DataFrame(duplicate_rows)
duplicate_rows
```

Out[22]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontag
4986	Residential Land(Land and Building)	Chuo Ward	2.0	230000000	1K	80	Rectangular Shaped	7
4994	Residential Land(Land and Building)	Chuo Ward	2.0	230000000	1K	80	Rectangular Shaped	7
9353	Residential Land(Land and Building)	Minato Ward	1.0	360000000	1K	160	Rectangular Shaped	8
9354	Residential Land(Land and Building)	Minato Ward	1.0	360000000	1K	160	Rectangular Shaped	8
14686	Residential Land(Land and Building)	Shinjuku Ward	5.0	59000000	1K	50	Rectangular Shaped	7
14690	Residential Land(Land and Building)	Shinjuku Ward	5.0	59000000	1K	50	Rectangular Shaped	7
18342	Residential Land(Land and Building)	Shinjuku Ward	4.0	400000000	1K	145	Semi-rectangular Shaped	9
18343	Residential Land(Land and Building)	Shinjuku Ward	4.0	400000000	1K	145	Semi-rectangular Shaped	9

Now I check the shape of the data before and after dropping the duplicates. Finally I confirm that the data doesnt have any duplicates.

In [23]:

```
print("before: ", df.shape)

df = df.drop_duplicates(keep='first')
print("after: ", df.shape)
```

before: (2765, 13)
after: (2761, 13)

In [24]:

```
duplicate_rows = df.duplicated()

if duplicate_rows.any():
    print("The DataFrame has duplicate rows.")
else:
    print("The DataFrame does not have duplicate rows.")
```

The DataFrame does not have duplicate rows.

6 - Checking the Numerical Variables

Now I want to check each numerical variable for any possible non-numerical values. This check is necessary as numeric values can be classified as an `int`, `float`, or even a `str`. The columns that are of interest are: `Area`, `Frontage`, `MinuteToClosestStation`, `BuildingYear`, `Price`. Since I only have five numerical variables to check, I will manually check each one

```
In [25]: print(df.Area.apply(type).value_counts())
print(df.Frontage.apply(type).value_counts())
print(df.MinuteToClosestStation.apply(type).value_counts())
print(df.BuildingYear.apply(type).value_counts())
print(df.Price.apply(type).value_counts())
```

```
<class 'int'>      2761
Name: Area, dtype: int64
<class 'float'>    2761
Name: Frontage, dtype: int64
<class 'float'>    2761
Name: MinuteToClosestStation, dtype: int64
<class 'float'>    2761
Name: BuildingYear, dtype: int64
<class 'int'>      2761
Name: Price, dtype: int64
```

Luckily, each numerical variable consists of either an `int` or `float`. Moreover, each variable has exactly 2761 values, which is the number of rows that our current dataframe, as seen below.

```
In [26]: df.shape
```

```
Out[26]: (2761, 13)
```

7 - Checking the Categorical Variables

Lastly, I will check the values of the categorical variable. These variables are: `Type`, `Municipality`, `FloorPlan`, `LandShape`, `Structure`, `Direction`, `Use`, `Renovation`. It is important to check the categorical variables as some may have multiple values according to the documentations found on the Kaggle page associated with the data found [here](#).

Type

```
In [27]: print(df['Type'].unique())
```

```
['Residential Land(Land and Building)']
```

```
In [28]: print(df['Type'].value_counts())
```

```
Residential Land(Land and Building)    2761
Name: Type, dtype: int64
```

Municipality


```
In [29]: print(df['Municipality'].unique())
```

```
['Chiyoda Ward' 'Chuo Ward' 'Minato Ward' 'Shinjuku Ward' 'Bunkyo Ward']
```

```
In [30]: print(df['Municipality'].value_counts())
```

```
Shinjuku Ward    1015
Minato Ward      660
Chuo Ward        581
Chiyoda Ward     434
Bunkyo Ward       71
Name: Municipality, dtype: int64
```

FloorPlan

```
In [31]: print(df['FloorPlan'].unique())
```

```
['1K']
```

```
In [32]: print(df['FloorPlan'].value_counts())
```

```
1K    2761
Name: FloorPlan, dtype: int64
```

LandShape

```
In [33]: print(df['LandShape'].unique())
```

```
['Trapezoidal Shaped' 'Rectangular Shaped' 'Semi-rectangular Shaped'
 'Semi-trapezoidal Shaped' 'Irregular Shaped' 'Semi-shaped'
 'Semi-square Shaped' 'Flag-shaped etc.' 'Square Shaped' 'Other LandShape']
```

```
In [34]: print(df['LandShape'].value_counts())
```

```
Semi-rectangular Shaped    1014
Rectangular Shaped         731
Irregular Shaped           423
Semi-trapezoidal Shaped    197
Semi-square Shaped         157
Semi-shaped                111
Trapezoidal Shaped         78
Flag-shaped etc.           28
Square Shaped              21
Other LandShape             1
Name: LandShape, dtype: int64
```

Structure

```
In [35]: print(df['Structure'].unique())
```

```
['S' 'RC' 'SRC' 'W' 'LS' 'S, LS' 'B' 'SRC, RC' 'S, W' 'RC, W' 'RC, S'
 'W, B']
```

It seems that some of the `Structure` values are a combination of two different Structure types (ex. `'S, LS'`). For simplicity my strategy will be to split these values up and only keep the first Structure type (ex. split `'S, LS'` and only keep `'S'`)

```
In [36]: df['Structure'] = df['Structure'].apply(lambda x: x.split(',')[0])
```

```
In [37]: print(df['Structure'].value_counts())
```

```
RC      1014
W       795
S       551
SRC     348
LS       50
B        3
Name: Structure, dtype: int64
```

Direction

```
In [38]: print(df['Direction'].unique())

['Northwest' 'Southwest' 'South' 'Northeast' 'East' 'Southeast' 'West'
 'North' 'No facing road']
```

```
In [39]: print(df['Direction'].value_counts())

Northwest      382
Southwest      373
Southeast      371
West           344
Northeast      335
South          329
East           307
North          304
No facing road   16
Name: Direction, dtype: int64
```

Use

```
In [40]: print(df['Use'].unique())
```

['Office, Shop' 'Housing Complex, Office' 'House, Office, Workshop'
'House, Office, Shop' 'Other' 'Office, Warehouse' 'Parking Lot'
'Office, Parking Lot' 'Housing Complex, Office, Shop' 'House' 'Office'
'Housing Complex' 'Housing Complex, Shop' 'House, Shop' 'Warehouse, Shop'
'Housing Complex, Office, Warehouse, Shop' 'House, Workshop, Shop'
'Warehouse' 'House, Office' 'Office, Warehouse, Parking Lot'
'House, Factory, Office, Shop' 'House, Office, Other' 'Shop'
'House, Office, Warehouse, Shop' 'Office, Other'
'Office, Warehouse, Shop' 'House, Warehouse, Shop, Other' 'Workshop'
'House, Office, Parking Lot' 'Office, Workshop, Warehouse' 'Factory'
'House, Warehouse' 'House, Housing Complex, Factory, Warehouse'
'Office, Parking Lot, Shop' 'House, Factory, Office'
'Housing Complex, Office, Warehouse, Parking Lot'
'House, Office, Warehouse' 'House, Workshop' 'House, Office, Shop, Other'
'House, Parking Lot' 'House, Parking Lot, Shop'
'Housing Complex, Parking Lot' 'Housing Complex, Warehouse'
'House, Factory' 'Warehouse, Parking Lot' 'Office, Warehouse, Other'
'Housing Complex, Office, Parking Lot, Shop'
'House, Office, Parking Lot, Shop' 'House, Other'
'House, Warehouse, Shop' 'Shop, Other' 'House, Workshop, Warehouse, Shop'
'House, Housing Complex, Office, Shop'
'Housing Complex, Office, Warehouse' 'Office, Shop, Other'
'House, Office, Warehouse, Parking Lot' 'Housing Complex, Other'
'Workshop, Warehouse, Parking Lot' 'Factory, Office, Warehouse'
'Factory, Office' 'House, Housing Complex, Factory, Office'
'House, Workshop, Warehouse' 'Office, Workshop, Shop'
'Housing Complex, Workshop' 'Office, Workshop' 'House, Shop, Other'
'House, Office, Workshop, Shop'
'House, Housing Complex, Office, Warehouse'
'Housing Complex, Office, Workshop' 'House, Housing Complex, Office'
'House, Warehouse, Other' 'House, Housing Complex, Shop'
'House, Housing Complex, Workshop'
'House, Housing Complex, Factory, Office, Workshop, Warehouse'
'Housing Complex, Factory, Office, Parking Lot'
'House, Housing Complex, Parking Lot'
'House, Warehouse, Parking Lot, Shop' 'Housing Complex, Factory, Office'
'Office, Parking Lot, Other' 'Housing Complex, Office, Parking Lot'
'Housing Complex, Parking Lot, Shop'
'Housing Complex, Workshop, Warehouse'
'Housing Complex, Office, Warehouse, Parking Lot, Shop'
'Housing Complex, Factory'
'Housing Complex, Office, Warehouse, Shop, Other'
'House, Housing Complex, Workshop, Shop'
'Housing Complex, Warehouse, Parking Lot' 'House, Housing Complex'
'Parking Lot, Shop' 'Housing Complex, Warehouse, Shop'
'House, Housing Complex, Office, Workshop' 'Parking Lot, Other'
'Workshop, Shop' 'House, Warehouse, Parking Lot'
'House, Office, Parking Lot, Shop, Other'
'House, Housing Complex, Parking Lot, Other'
'House, Parking Lot, Shop, Other' 'Office, Warehouse, Parking Lot, Shop'
'House, Office, Parking Lot, Other' 'Factory, Workshop'
'Housing Complex, Office, Other'
'House, Office, Warehouse, Parking Lot, Shop' 'Workshop, Warehouse'
'Housing Complex, Factory, Shop'
'House, Workshop, Warehouse, Parking Lot'
'House, Office, Warehouse, Shop, Other' 'Factory, Office, Other'
'Office, Workshop, Warehouse, Parking Lot'
'Housing Complex, Workshop, Shop' 'Office, Parking Lot, Shop, Other'
'House, Housing Complex, Office, Warehouse, Parking Lot']

Like with the `Structure` column, `Use` has entries that have a combination of values. Again, I will split them up and only consider the first `Use` value.

```
In [41]: df['Use'] = df['Use'].apply(lambda x: x.split(',')[0])
```

```
In [42]: print(df['Use'].value_counts())
```

```
House          1300
Office          705
Housing Complex 559
Shop            117
Other           43
Warehouse       18
Factory         9
Parking Lot     6
Workshop        4
Name: Use, dtype: int64
```

Renovation

```
In [43]: print(df['Renovation'].unique())
```

```
['Not yet']
```

```
In [44]: print(df['Renovation'].value_counts())
```

```
Not yet    2761
Name: Renovation, dtype: int64
```

Step 2 Exploratory Data Analysis (EDA)

Now that the Data is properly cleaned I can conduct Exploratory Data Analysis. Once this is completed I can move to Correlation analysis and building a ML model (which will be done in a separate Notebook).

1 - Discriptive Statical Analysis with SweetViz package

I recently discovered the SweetViz package. SweetViz is an open-source Python library that generates beautiful, high-density visualizations to kickstart EDA (Exploratory Data Analysis) with just two lines of code. Output is a fully self-contained HTML application or can be embedded within a Jupyter Notebook.

I opt to include the SweetViz visualization within the notebook. The great thing about SweetViz output is that it is fully interactive. Each tab contains basic statistical analysis of a Variable from the DataFrame. The `ASSOCIATIONS` button will provide a correlation matrix.

```
In [45]: analyze_report = sv.analyze(df)
analyze_report.show_notebook()
```

| [0%] 00:00 -> (? left)

1

Type

VALUES: 2,761 (100%)
 MISSING: ---
 DISTINCT: 1 (<1%)

Residential Land

2

Municipality

VALUES: 2,761 (100%)
 MISSING: ---
 DISTINCT: 5 (<1%)

3

MinuteToClosestStation

VALUES:	2,761 (100%)	MAX	16.0	RANGE	16.0
MISSING:	---	95%	9.0	IQR	3.00
DISTINCT:	17 (<1%)	Q3	6.0	STD	2.58
		AVG	4.6	VAR	6.68
		MEDIAN	4.0		
ZEROES:	10 (<1%)	Q1	3.0	KURT.	0.076
		5%	1.0	SKEW	0.724
		MIN	0.0	SUM	12,597

4

Price

2 - Focusing in on Numerical Variables

I feel it is important to focus in on the Numerical variables a bit more than what SweetViz has outputted. Specifically I will create a KDE plot for each Numerical Variables as well as provide a written interpretation for each Numerical Variables. Below I will generate a KDE for the Numerical variables as well as provide some interpretations using the SweetViz output.

First I create a function to plot the KDE along with the Mean in black and median in blue

```
In [46]: def kde_plot(var):
```

```

plt.figure(figsize = (8,5))
sns.histplot(df[var], kde = True, color= 'green')
sns.despine(left = True)


mean = df[var].mean()
median = df[var].median()

plt.axvline(mean, color = 'black', linestyle = 'solid')
plt.axvline(median, color = 'blue', linestyle = 'solid')
plt.xlabel('')
plt.ylabel('')
plt.title("KDE plot of {}".format(var))
plt.legend

return plt.show()

```

In [47]: `kde_plot('Area')`




Notes on Area:

- The average Area is 170 square meters
- Maximum: 2000 square meters
- Minimum: 30 square meters
- Half of the real estate are less than 105 square meters
- 25% of the real estate have an area less than 65 square meters
- 75% of the real estate have an area greater than 190 square meters
- 50% of the real estate have an area between 65 and 190 square meters
- The KDE clearly shows that Area is skewed right

In [48]: `kde_plot('Frontage')`

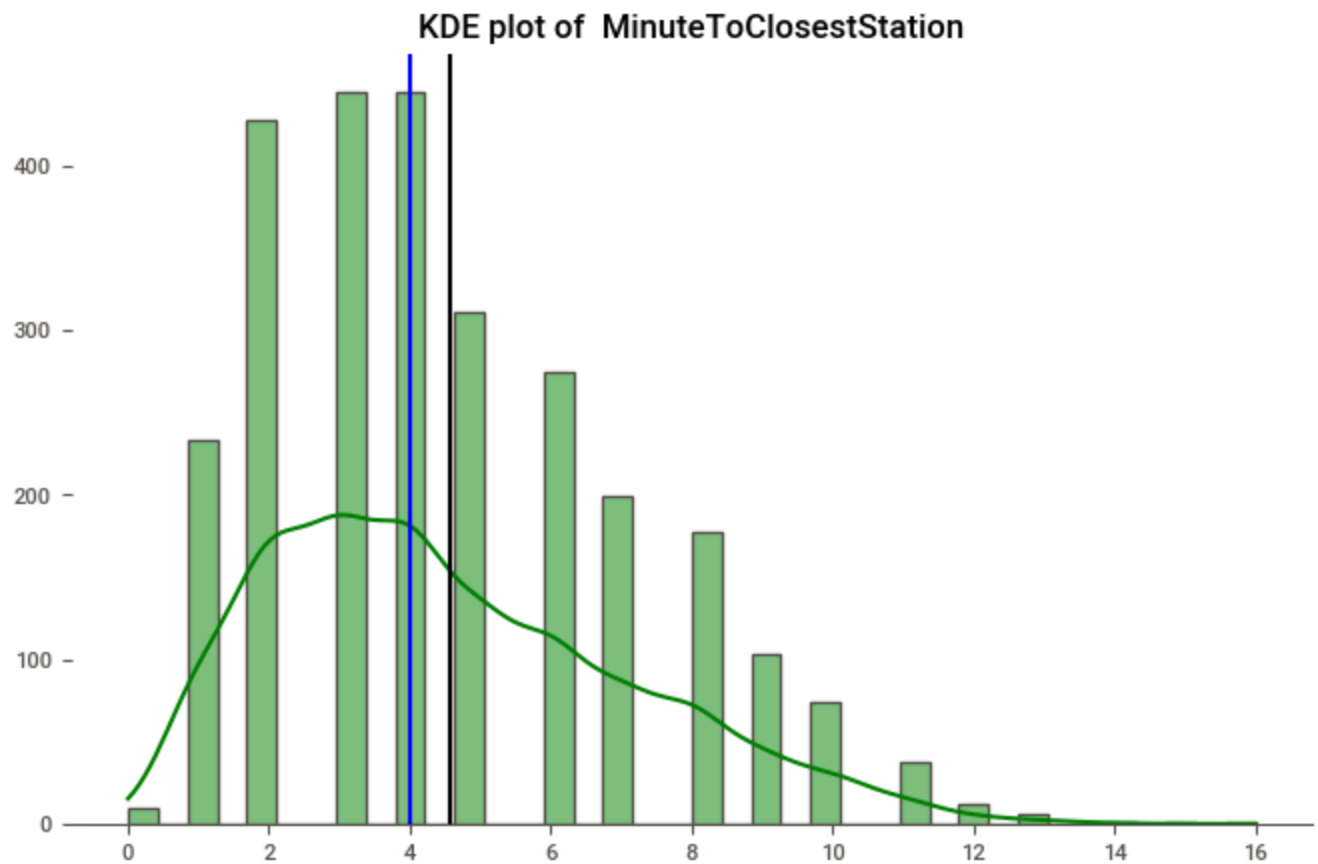
KDE plot of Frontage



Notes on Frontage:

- The average Frontage is 9.6 meters
- Maximum: 50 meters
- Minimum: 1 meters
- Half of the real estate have Frontage that is less than 8 meters
- 25% of the real estate have Frontage that is less than 5.5 meters
- 75% of the real estate have Frontage that is greater than 11.5 meters
- 50% of the real estate have Frontage that is between 5.5 and 11.5 meters
- The KDE clearly shows that Frontage is skewed right


```
In [49]: kde_plot('MinuteToClosestStation')
```



Notes on MinuteToClosestStation:

- The average time to nearest station is 4.6 minutes
- Maximum: 16 minutes to nearest station
- Minimum: less than 1 minute to nearest station
- Half of the real estate is less than 4 minutes to the nearest station
- 25% of the real estate is less than 3 minutes to the nearest station
- 75% of the real estate is greater than 6 minutes to the nearest station
- 50% of the real estate is between 3 and 6 minutes to the nearest station
- The KDE shows that the time to the nearest station resembles a normal distribution but is skewed right

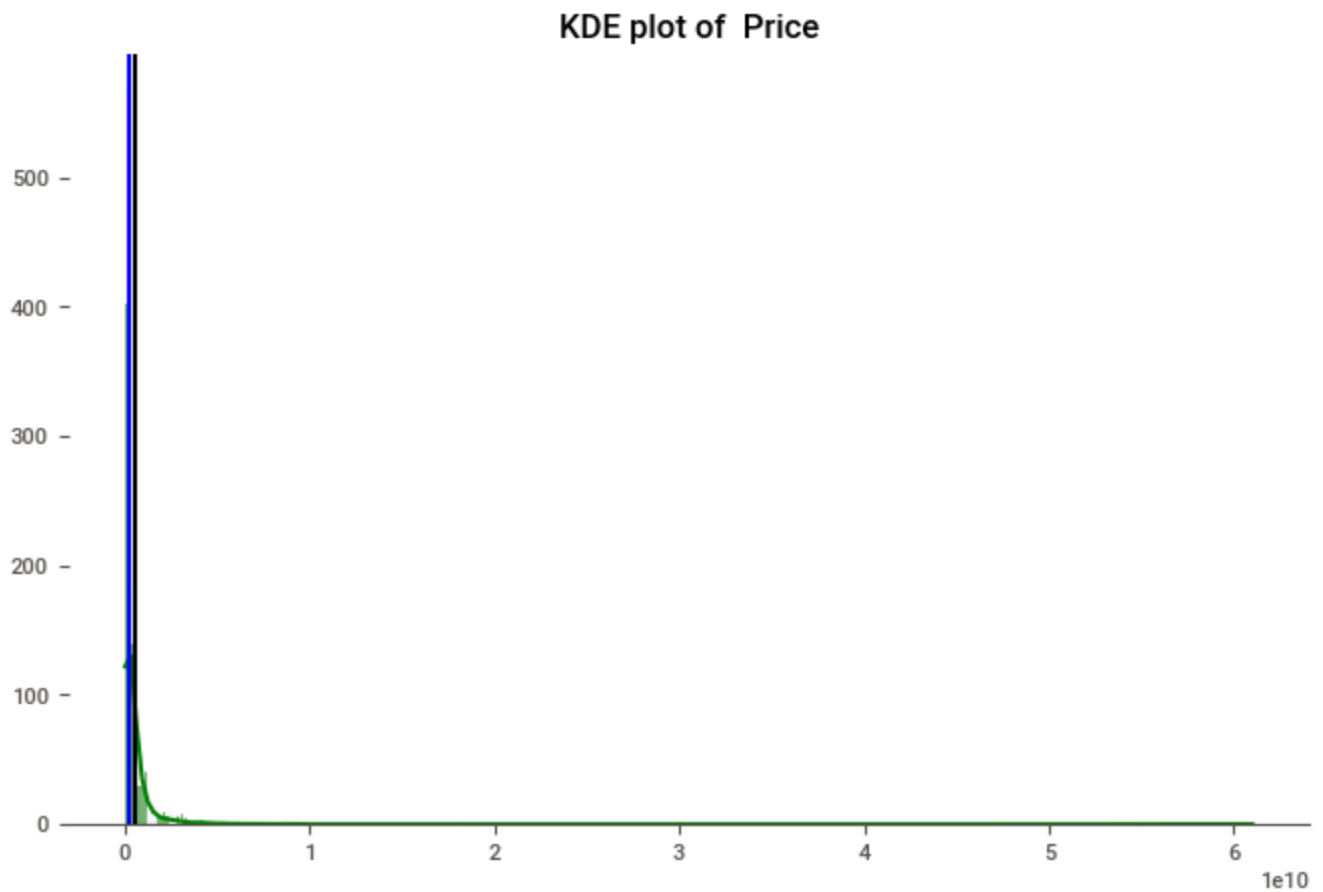
```
In [50]: kde_plot('BuildingYear')
```

Notes on BuildingYear:

- The average year to build is 1987
- Maximum: The newest house was built in 2020
- Minimum: The oldest house was built in 1945
- Half of all the real estate was built before 1988
- 25% of all the real estate was built before or on 1974
- 75% of all the real estate was built before or on 2000
- 50% of all the real estate was built between 1974 and 2000
- The KDE shows that the distribution of building year is somewhere between a normal and bimodal distribution

```
In [51]: kde_plot('Price')
```



Notes on Price:

- The average price was 0.5 billion Yen
- Maximum: 61 Billion Yen
- Minimum: 0 Billion Yen
- Half of all the real estate is worth less than 0.2 billion Yen
- 25% of all the real estate is worth less than 0.1 billion Yen
- 75% of all the real estate is worth less than 0.5 billion Yen
- 50% of all the real estate is worth between 0.1 and 0.5
- The KDE shows that the distribution of price is extremely skewed right

Now I will present the real estate with the maximum and minimums of each numerical value. This is just so we can conduct some more of a visual analysis of the properties with maximum and minimum numerical traits.

Minimum Area

```
In [52]: df[df['Area'] == df['Area'].min()]
```

Out[52]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontag
419	Residential Land(Land and Building)	Chiyoda Ward	2.0	17000000	1K	30	Rectangular Shaped	4
2461	Residential Land(Land and Building)	Chuo Ward	1.0	320000000	1K	30	Rectangular Shaped	3
2499	Residential Land(Land and Building)	Chuo Ward	2.0	190000000	1K	30	Semi-rectangular Shaped	4
2525	Residential Land(Land and Building)	Chuo Ward	2.0	120000000	1K	30	Trapezoidal Shaped	5
3387	Residential Land(Land and Building)	Chuo Ward	2.0	50000000	1K	30	Rectangular Shaped	3
3388	Residential Land(Land and Building)	Chuo Ward	2.0	51000000	1K	30	Rectangular Shaped	3
5213	Residential Land(Land and Building)	Chuo Ward	3.0	150000000	1K	30	Semi-rectangular Shaped	3
5832	Residential Land(Land and Building)	Chuo Ward	3.0	36000000	1K	30	Rectangular Shaped	3
8425	Residential Land(Land and Building)	Minato Ward	8.0	43000000	1K	30	Semi-shaped	5
9040	Residential Land(Land and Building)	Minato Ward	4.0	27000000	1K	30	Semi-rectangular Shaped	7
9714	Residential Land(Land and Building)	Minato Ward	6.0	23000000	1K	30	Semi-rectangular Shaped	4
10367	Residential Land(Land and Building)	Minato Ward	4.0	24000000	1K	30	Semi-rectangular Shaped	3
11721	Residential Land(Land	Chuo Ward	4.0	9000000	1K	30	Semi-rectangular	4

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage
	and Building)						Shaped	
11993	Residential Land(Land and Building)	Shinjuku Ward	5.0	4500000	1K	30	Semi-square Shaped	3
12579	Residential Land(Land and Building)	Shinjuku Ward	4.0	17000000	1K	30	Rectangular Shaped	3
13260	Residential Land(Land and Building)	Shinjuku Ward	5.0	47000000	1K	30	Semi-trapezoidal Shaped	4
14636	Residential Land(Land and Building)	Chuo Ward	4.0	33000000	1K	30	Semi-rectangular Shaped	7
15144	Residential Land(Land and Building)	Shinjuku Ward	3.0	22000000	1K	30	Semi-rectangular Shaped	4
16219	Residential Land(Land and Building)	Minato Ward	2.0	40000000	1K	30	Irregular Shaped	3
16628	Residential Land(Land and Building)	Shinjuku Ward	5.0	20000000	1K	30	Rectangular Shaped	3
17544	Residential Land(Land and Building)	Shinjuku Ward	4.0	20000000	1K	30	Semi-rectangular Shaped	4
18388	Residential Land(Land and Building)	Chuo Ward	2.0	170000000	1K	30	Rectangular Shaped	3
18406	Residential Land(Land and Building)	Shinjuku Ward	0.0	140000000	1K	30	Irregular Shaped	7
18527	Residential Land(Land and Building)	Shinjuku Ward	8.0	21000000	1K	30	Semi-shaped	6
19082	Residential Land(Land and Building)	Bunkyo Ward	5.0	39000000	1K	30	Irregular Shaped	7

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage
19750	Residential Land(Land and Building)	Bunkyo Ward	5.0	37000000	1K	30	Rectangular Shaped	4
19988	Residential Land(Land and Building)	Bunkyo Ward	6.0	40000000	1K	30	Rectangular Shaped	3
19989	Residential Land(Land and Building)	Bunkyo Ward	6.0	36000000	1K	30	Rectangular Shaped	3

Maximum Area

```
In [53]: df[df['Area'] == df['Area'].max()]
```

Out[53]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Front
1053	Residential Land(Land and Building)	Chiyoda Ward	6.0	7000000000	1K	2000	Rectangular Shaped	
6465	Residential Land(Land and Building)	Minato Ward	1.0	30000000000	1K	2000	Irregular Shaped	
6662	Residential Land(Land and Building)	Minato Ward	7.0	7500000000	1K	2000	Irregular Shaped	
9930	Residential Land(Land and Building)	Minato Ward	3.0	6600000000	1K	2000	Semi-rectangular Shaped	
11060	Residential Land(Land and Building)	Minato Ward	9.0	4700000000	1K	2000	Irregular Shaped	
15053	Residential Land(Land and Building)	Minato Ward	6.0	7000000000	1K	2000	Semi-square Shaped	
15142	Residential Land(Land and Building)	Minato Ward	1.0	4000000000	1K	2000	Semi-rectangular Shaped	
15812	Residential Land(Land and Building)	Chuo Ward	7.0	14000000000	1K	2000	Semi-shaped	
17597	Residential Land(Land and Building)	Shinjuku Ward	5.0	3900000000	1K	2000	Semi-shaped	
17778	Residential Land(Land and Building)	Shinjuku Ward	9.0	4400000000	1K	2000	Irregular Shaped	
18711	Residential Land(Land and Building)	Minato Ward	4.0	61000000000	1K	2000	Irregular Shaped	
<div><div></div><div></div></div>								

Minimum Frontage

In [54]: df[df['Frontage'] == df['Frontage'].min()]

Out[54]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage
17351	Residential Land(Land and Building)	Shinjuku Ward	8.0	7600000	1K	40	Flag-shaped etc.	1.0

Maximum Frontage

In [55]:

```
df[df['Frontage'] == df['Frontage'].max()]
```

Out[55]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Front
6465	Residential Land(Land and Building)	Minato Ward	1.0	30000000000	1K	2000	Irregular Shaped	1.0
9930	Residential Land(Land and Building)	Minato Ward	3.0	6600000000	1K	2000	Semi-rectangular Shaped	1.0
15812	Residential Land(Land and Building)	Chuo Ward	7.0	14000000000	1K	2000	Semi-shaped	1.0
16263	Residential Land(Land and Building)	Minato Ward	3.0	2700000000	1K	1300	Semi-rectangular Shaped	1.0
16268	Residential Land(Land and Building)	Shinjuku Ward	6.0	1900000000	1K	1400	Semi-rectangular Shaped	1.0
17597	Residential Land(Land and Building)	Shinjuku Ward	5.0	3900000000	1K	2000	Semi-shaped	1.0
17778	Residential Land(Land and Building)	Shinjuku Ward	9.0	4400000000	1K	2000	Irregular Shaped	1.0
18711	Residential Land(Land and Building)	Minato Ward	4.0	61000000000	1K	2000	Irregular Shaped	1.0

Minimum Time to Nearest Station

In [56]:

```
df[df['MinuteToClosestStation'] == df['MinuteToClosestStation'].min()]
```

Out[56]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Fronta
3722	Residential Land(Land and Building)	Chuo Ward	0.0	3800000000	1K	390	Semi-rectangular Shaped	1
4961	Residential Land(Land and Building)	Chuo Ward	0.0	140000000	1K	50	Semi-rectangular Shaped	
5181	Residential Land(Land and Building)	Chuo Ward	0.0	2500000000	1K	470	Irregular Shaped	1
5259	Residential Land(Land and Building)	Chuo Ward	0.0	58000000	1K	60	Rectangular Shaped	
5790	Residential Land(Land and Building)	Chuo Ward	0.0	110000000	1K	45	Rectangular Shaped	
13500	Residential Land(Land and Building)	Chuo Ward	0.0	200000000	1K	55	Semi-rectangular Shaped	
14363	Residential Land(Land and Building)	Shinjuku Ward	0.0	2300000000	1K	240	Semi-trapezoidal Shaped	1
18406	Residential Land(Land and Building)	Shinjuku Ward	0.0	140000000	1K	30	Irregular Shaped	
18614	Residential Land(Land and Building)	Minato Ward	0.0	2300000000	1K	300	Rectangular Shaped	1
19603	Residential Land(Land and Building)	Minato Ward	0.0	520000000	1K	75	Trapezoidal Shaped	

Maximum Time to Nearest Station

In [57]: `df[df['MinuteToClosestStation'] == df['MinuteToClosestStation'].max()]`

Out[57]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontag
19975	Residential Land(Land and Building)	Minato Ward	16.0	110000000	1K	85	Semi- square Shaped	8.

◀

▶

Minimum Building Year

In [58]:


df[df['BuildingYear'] == df['BuildingYear'].min()]

Maximum Building Year

```
In [59]: df[df['BuildingYear'] == df['BuildingYear'].max()]
```

Out[59]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage
17895	Residential Land(Land and Building)	Shinjuku Ward	5.0	76000000	1K	75	Flag-shaped etc.	2.5




Minimum Price

```
In [60]: df[df['Price'] == df['Price'].min()]
```

Out[60]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage	B
7392	Residential Land(Land and Building)	Minato Ward	4.0	520000	1K	45	Semi-rectangular Shaped	5.0	




Maximum Price

```
In [61]: df[df['Price'] == df['Price'].max()]
```

Out[61]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Front
18711	Residential Land(Land and Building)	Minato Ward	4.0	61000000000	1K	2000	Irregular Shaped	



3 - Outliers

In the previous sections, we saw some variables have outliers, and before wrapping up the EDA I want to address this issue. The biggest outliers are seen in `Price` and looking at the correlation matrix in the SweetViz dashboard I see that `Price` is correlated to `Area`. So, with some feature engineering I will handle these outliers. To start, I will first create a new feature Price per Area: `PPA`

```
In [62]: df['PPA'] = df['Price'] / df['Area']  
df.head(2)
```

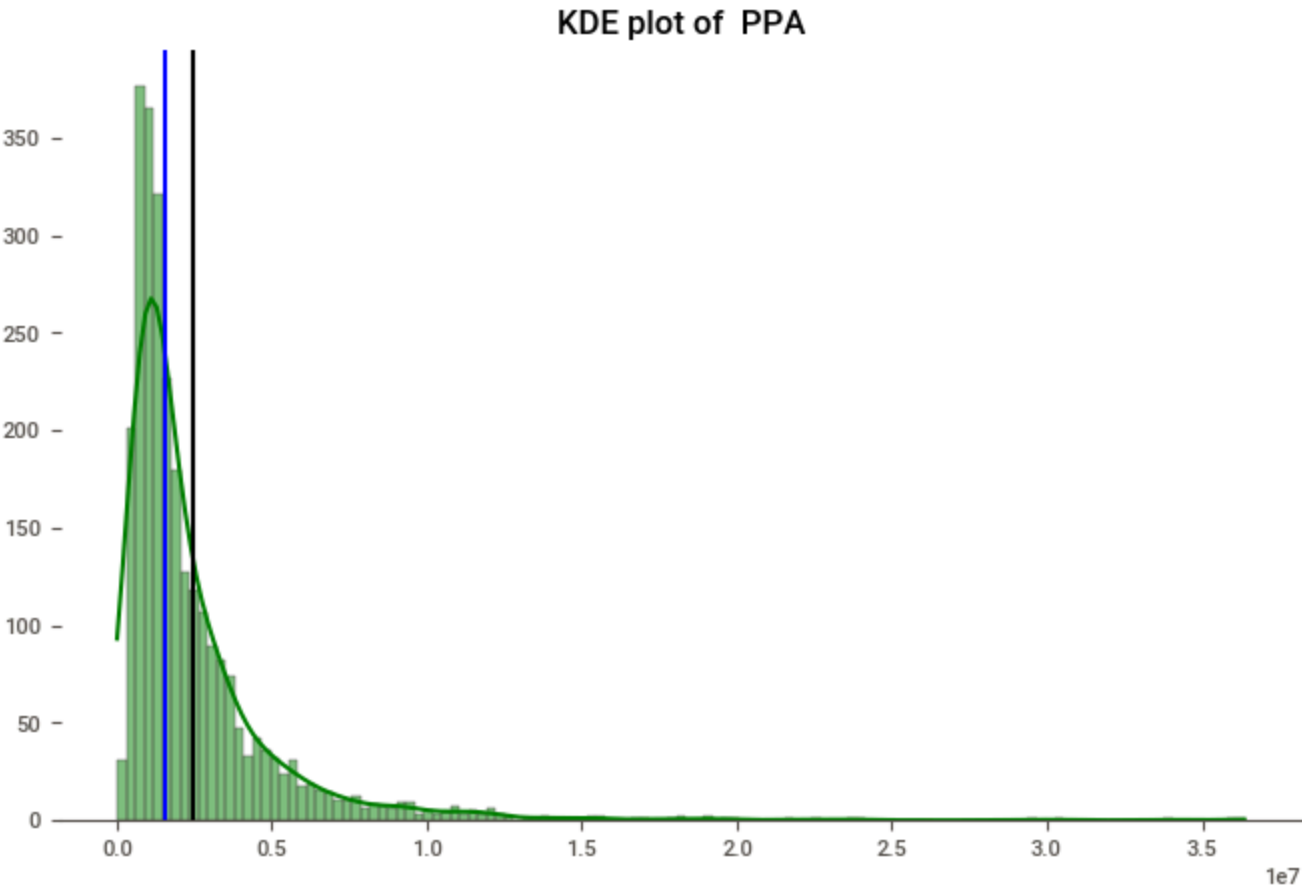
Out[62]:

	Type	Municipality	MinuteToClosestStation	Price	FloorPlan	Area	LandShape	Frontage	B
2	Residential Land(Land and Building)	Chiyoda Ward	2.0	400000000	1K	110	Trapezoidal Shaped	9.0	
3	Residential Land(Land and Building)	Chiyoda Ward	1.0	180000000	1K	50	Rectangular Shaped	5.2	

Now the KDE of PPA

In [63]:

```
kde_plot('PPA')
```



The KDE shows that PPA is extremely right skewed and contains very large outlier. Therefore we cannot methods like z-score, as PPA does not behave normally. Instead, I will use the Median Absolute Deviation (MAD) technique. This is a robust method for distribution with a heavy outlier effect. More information on the MAD can be found [here](#). I will calculate the MAD of PPA using the `scipy.stats` library below.

In [64]:

```
median_abs_deviation(df["PPA"])
```

Out[64]: 788888.8888888889

Now, I filter dataset (removing outlier) based on the MAD technique.

In [65]:

```
median_prices = {"Tokyo": median_abs_deviation(df["PPA"])}
```

```

#call province median from the dictionary
Median = median_prices["Tokyo"]

#difference between each price_per_area with the called median
df['Median_Diff'] = 0

for index, row in df.iterrows():
    median_diff = abs(row['PPA'] - Median)
    df.at[index, 'Median_Diff'] = median_diff

#calculate the median of new column
MAD = df['Median_Diff'].median()

#determine treshold
threshold = MAD * 3

#detect and filter rows based on outlier
df = df[~(df['Median_Diff'] > threshold)]

#remove the differenece column
df = df.drop(['Median_Diff'], axis=1)

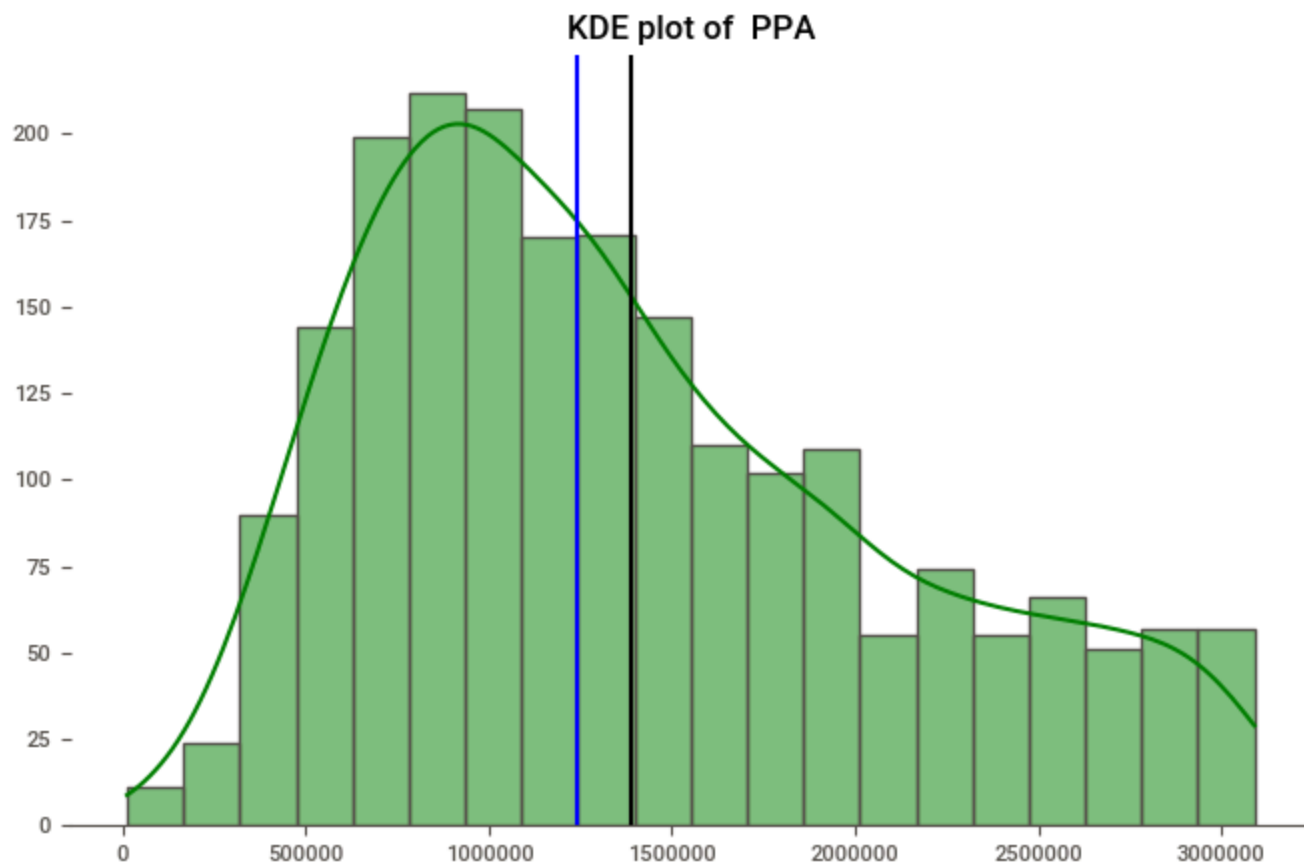
```

In [66]: `df.shape`

Out[66]: (2111, 14)

Note that the shape of the Dataframe has been reduced, an indication that outliers have been removed. We can check the KDE of `PPA` once again and visually confirm the removal of extreme outliers.

In [67]: `kde_plot('PPA')`



Now that we removed outliers based on `PPA` we can remove the `PPA` column, reset the index, and proceed to save our processed dataframe using the `pickle` package.

```
In [68]: df=df.drop(['PPA'], axis=1)
```

```
In [69]: df = df[df.Use != 'Parking Lot']
```

4 - Saving Data using `pickle`

This concludes the first part of this project. Using the `pickle` package I will save our current dataframe. In the next part I will continue this project.

```
In [929... with open('TokyoSub_Final.pickle', 'wb') as file:
            pickle.dump(df, file)
```
